

DTSA 5509 Final

June 21, 2022

0.1 Section 1: Project Topic

Boston House Prices

My general goal for this project is to explore the different variables that go into determining median house values, and see if I can use different machine learning models to predict these values.

0.2 Section 2: Data Source

<https://www.kaggle.com/datasets/fedesoriano/the-boston-houseprice-data>

The original source of the dataset is StatLib - Carnegie Mellon University

This dataset is derived from The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978.

0.3 Section 3: Data Description

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import scipy as sp
import scipy.stats as stats
import statsmodels.formula.api as smf
import statsmodels.api as sm
from sklearn import metrics
from sklearn.model_selection import cross_val_score
import sklearn
from sklearn.metrics import mean_squared_error
from sklearn.svm import SVR
from sklearn.svm import LinearSVC
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestRegressor
```

```

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.base import clone
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import warnings
import datetime
warnings.filterwarnings("ignore")
%matplotlib inline

```

```

[2]: data = pd.read_csv("boston.csv")
data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   CRIM        506 non-null   float64
 1   ZN          506 non-null   float64
 2   INDUS       506 non-null   float64
 3   CHAS        506 non-null   int64
 4   NOX         506 non-null   float64
 5   RM          506 non-null   float64
 6   AGE         506 non-null   float64
 7   DIS         506 non-null   float64
 8   RAD         506 non-null   int64
 9   TAX         506 non-null   float64
10  PTRATIO     506 non-null   float64
11  B           506 non-null   float64
12  LSTAT       506 non-null   float64
13  MEDV        506 non-null   float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB

```

Description of all variables:

- 1) CRIM: per capita crime rate by town
- 2) ZN: proportion of residential land zoned for lots over 25,000 sq.ft.
- 3) INDUS: proportion of non-retail business acres per town
- 4) CHAS: Charles River dummy variable (1 if tract bounds river; 0 otherwise)
- 5) NOX: nitric oxides concentration (parts per 10 million) [parts/10M]
- 6) RM: average number of rooms per dwelling

- 7) AGE: proportion of owner-occupied units built prior to 1940
- 8) DIS: weighted distances to five Boston employment centres
- 9) RAD: index of accessibility to radial highways
- 10) TAX: full-value property-tax rate per 10,000[/10k]
- 11) PTRATIO: pupil-teacher ratio by town
- 12) B: The result of the equation $B=1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- 13) LSTAT: % lower status of the population
- 14) MEDV: Median value of owner-occupied homes in 1000 s[k]

```
[3]: print(data.shape)

print('Number of rows: ', data.shape[0])
print('Number of independent Variables: ', data.shape[1])
print('Target variable: MEDV')
```

```
(506, 14)
Number of rows: 506
Number of independent Variables: 14
Target variable: MEDV
```

```
[4]: data.describe()
```

```
[4]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677082	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	MEDV
--	-------	------

count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[ ]: ## Section 4: Data Cleaning
```

```
[5]: data.isnull().sum()
```

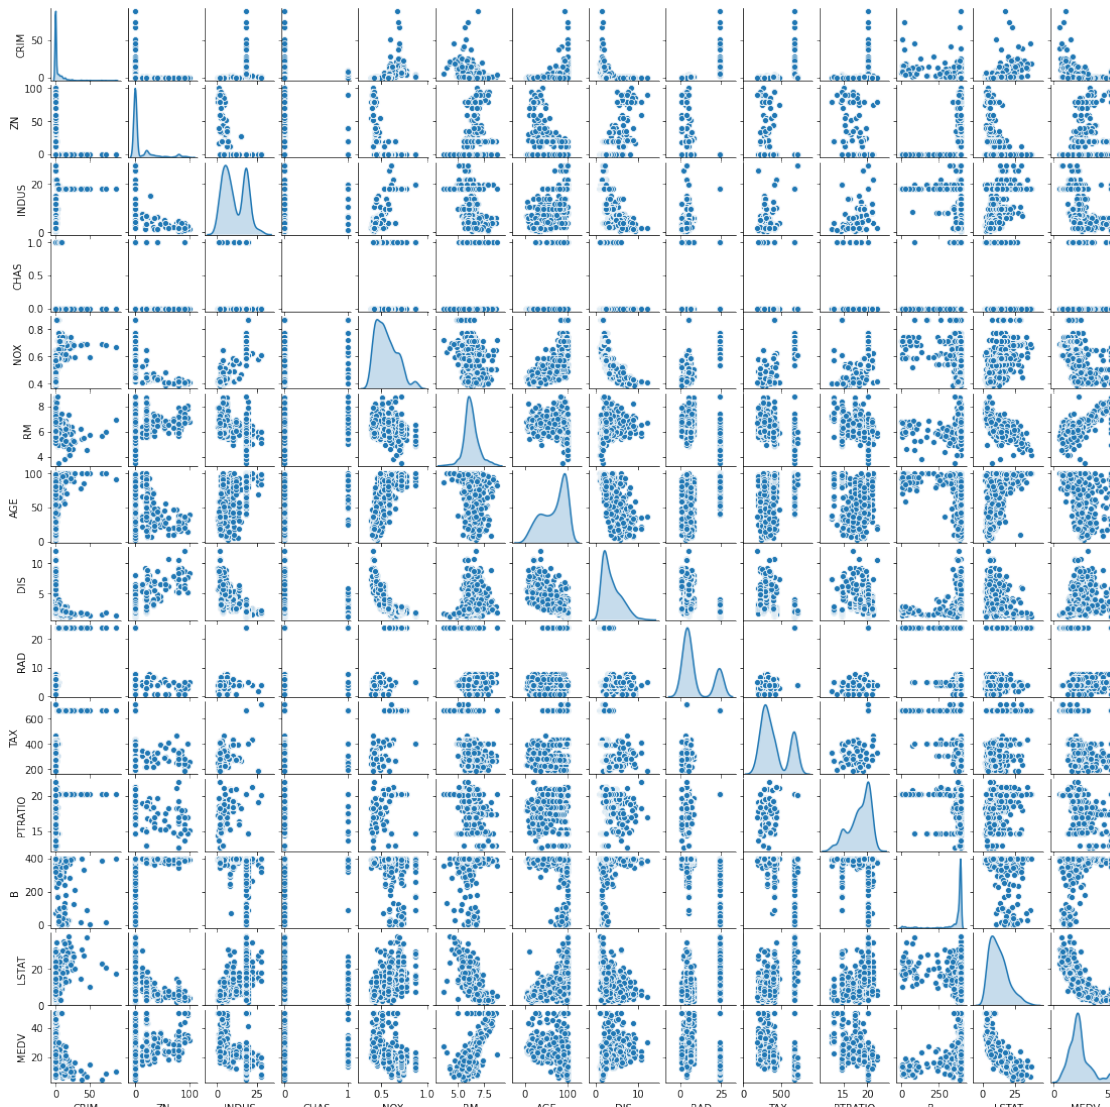
```
[5]: CRIM      0
      ZN       0
      INDUS   0
      CHAS    0
      NOX     0
      RM      0
      AGE     0
      DIS     0
      RAD     0
      TAX     0
      PTRATIO 0
      B       0
      LSTAT   0
      MEDV    0
      dtype: int64
```

Since there are no null values present, and I decided not to drop any of the columns from the dataset, I went ahead and proceeded straight away with my exploratory data analysis.

0.4 Section 5: Exploratory Data Analysis(EDA)


First, in order to see the relationships between all the variables, I decided to make a pair plot:

```
[6]: x=sns.pairplot(data, diag_kind="kde")
      x.fig.set_size_inches(15,15)
```



```
[7]: plt.figure(figsize=(16,6))
     sns.heatmap(data.corr(), annot = True)
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7f52f2d8b6d0>
```




From the matrix above, we can see that there is a strong correlation between the “RAD” and “TAX” variables.

Based off of this, I wanted to explore the relationship between these two variables, and create different visualizations to see the relationship between other variables as well.


```
[8]: sns.lmplot(x='TAX', y='RAD', data=data, aspect=2)
```

```
[8]: <seaborn.axisgrid.FacetGrid at 0x7f52ea3743d0>
```




```
[9]: sns.lmplot(x='RM', y='MEDV', data=data, aspect=2)
```

```
[9]: <seaborn.axisgrid.FacetGrid at 0x7f52ea8a4b50>
```




```
[10]: sns.violinplot(x='RM', y='MEDV', data=data, aspect=2)
plt.figure(figsize=(16,6))
```

[10]: <Figure size 1152x432 with 0 Axes>



<Figure size 1152x432 with 0 Axes>

```
[11]: sns.set_palette("gist_rainbow_r")
sns.jointplot(x="RM", y="MEDV", kind="hex", data=data )
plt.show()
```




I used several different visualizations to show the positive correlation between avg # of rooms per dwelling and median home value, a relationship which is more obvious to predict. I thought this would be the same case for the next relationship I explored. In this case, I was surprised to see the negative correlation between accessibility to radial highways and median house-value. I would have assumed that immediate access to highways would mean easier accesibility, and would increase the value of a home:

```
[12]: sns.lmplot(x='RAD', y='MEDV', data=data, aspect=2)
```



```
[12]: <seaborn.axisgrid.FacetGrid at 0x7f52df80ab10>
```



```
[ ]: ## Section 6: Build a Model
```

Given that MEDV is the output variable, I am going to split the dataset into training and test data sets based off of this variable:

```
[13]: X = data.drop('MEDV', axis=1)
      y = data['MEDV']
```

```
[14]: #Create test and train data
      X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.
      ↪8, random_state=42)
```

After splitting the data, I needed to decide which models I wanted to use.

The first model I decided to use is a Linear Regression Model. I chose this model, because I thought it would be the simplest model to start with. Before using that model, however, I wrote a small function that would calculate some important outputs:


```
[44]: def calculations(true, predicted):
      mae = metrics.mean_absolute_error(true, predicted)
      mse = metrics.mean_squared_error(true, predicted)
      rmse = np.sqrt(metrics.mean_squared_error(true, predicted))
      r2_sq = metrics.r2_score(true, predicted)
      print('MAE:', mae)
      print('MSE:', mse)
      print('RMSE:', rmse)
      print('R2 Square', r2_sq)
```

```
[45]: lin = LinearRegression(normalize=True)
lin.fit(X_train,y_train)
test_pred = lin.predict(X_test)
train_pred = lin.predict(X_train)

print('Test set')
calculations(y_test, test_pred)
print('Train set')
calculations(y_train, train_pred)
```

Test set
MAE: 3.1890919658878474
MSE: 24.29111947497351
RMSE: 4.928602182665336
R2 Square 0.6687594935356321
Train set
MAE: 3.31477162678323
MSE: 21.641412753226312
RMSE: 4.6520331848801675
R2 Square 0.7508856358979673

```
[47]: #Plotting prediction
pred = lin.predict(X_test)
plt.scatter(y_test, pred)
plt.show()
```




Originally, I was going to choose a DecisionTreeClassifier as my second model. However, I thought that this would work better on a binary variable, and it might take a long time to load. The second model I decided to use is a Random Forest Regressor:

```
[41]: randfor = RandomForestRegressor()
randfor.fit(X_train,y_train)
test_pred = randfor.predict(X_test)
train_pred = randfor.predict(X_train)

print('Test set')
calculations(y_test, test_pred)
print('Train set')
calculations(y_train, train_pred)
```

```
Test set
MAE: 2.0455
MSE: 8.528578186274512
RMSE: 2.9203729532843083
R2 Square 0.8837019199237375
Train set
MAE: 0.8702549504950492
MSE: 1.8557954480198025
RMSE: 1.3622758340438264
R2 Square 0.9786379332898228
```

```
[46]: #Plotting prediction
pred_1 = randfor.predict(X_test)
plt.scatter(y_test, pred_1)
plt.show()
```



The third model I chose is the Support Vector Algorithm:

```
[48]: svm= SVR(kernel='rbf', C=1000000, epsilon=0.001)
      svm.fit(X_train, y_train)

      test_pred = svm.predict(X_test)
      train_pred = svm.predict(X_train)

      print('Testing Values:')
      one=calculations(y_test, test_pred)
      print('-----')
      print('Training values')
      two=calculations(y_train, train_pred)
```

Testing Values:

MAE: 2.406948541994415

MSE: 13.177320758418691

RMSE: 3.6300579552424077

R2 Square 0.8203103646022145

Training values

MAE: 1.6575729516899593

MSE: 7.504002027415164

RMSE: 2.739343356977209

R2 Square 0.9136214111991737

```
[29]: ## Section 7: Discussion/Conclusion
```

- 1) All three models, although had high R2 values, were all examples of overfitting: testing MSE much higher than training data MSE
- 2) If I had to choose the best model, based off of R2 alone, I would choose the Random Forest Regressor
- 3) Had I done this project over, I definitely would have chosen to explore the CRIM- crime variable instead
- 4) Second, I would have chosen to split up the dataset into training and test data based off of a different variable, such as Crime, or even nitric-oxide concentration, just to see any change in results
- 5) Third, I would have chosen to focus my analysis more on the Charles River dummy variable, which would have allowed me to use different ML models
- 6) I was not fully happy with all of my choices, but I think I found some interesting insights